

#BOSTON HOUSING DATASET There are 506 samples and 14 feature variables in this dataset. The objective is to predict the value of prices of the house using the given features.

```
In [5]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
#Step 2: Import the Boston Housing dataset
from sklearn.datasets import load_boston
boston = load_boston()
```

```
In [8]: df= pd.DataFrame(boston.data)
```

```
In [9]: df
```

Out[9]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33
...	...	...	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.67
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.08
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.64
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.48
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.88

506 rows × 13 columns

```
In [11]: df.shape
```

Out[11]: (506, 13)

```
In [13]: df.columns= boston.feature_names
```

```
In [14]: df
```

Out[14]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33
...	...	...	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.67
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.08
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.64
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.48
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.88

506 rows × 13 columns

```
In [15]: df.head()
```

Out[15]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [17]: df['PRICE'] = boston.target
```

```
In [20]: df.isnull().sum()
```

Out[20]:

```
CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     0
```

```
In [25]: x = df.drop(['PRICE'], axis =1)
        y = df['PRICE']
```

```
In [27]: from sklearn.model selection import train test split
```

```
In [28]: xtrain, xtest, ytrain, ytest =train_test_split(x, y, test_size =0.2,random_state =0)
```

```
In [31]: import sklearn
        from sklearn.linear_model import LinearRegression
        lm = LinearRegression()
        model=lm.fit(xtrain, ytrain)
```

```
In [32]: lm.intercept
```

Out[32]: 38.09169492630278

```
In [33]: lm.coef
```

Out[33]: array([-1.19443447e-01, 4.47799511e-02, 5.48526168e-03, 2.34080361e+00,
 -1.61236043e+01, 3.70870901e+00, -3.12108178e-03, -1.38639737e+00,
 2.44178327e-01, -1.09896366e-02, -1.04592119e+00, 8.11010693e-03,
 -4.92792725e-01])

```
In [34]: ytrain_pred = lm.predict(xtrain)
        ytest_pred = lm.predict(xtest)
```


```
In [35]: df=pd.DataFrame(ytrain_pred,ytrain)
        df=pd.DataFrame(ytest_pred,ytest)
```

```
In [38]: from sklearn.metrics import mean_squared_error, r2_score
        mse = mean_squared_error(ytest, ytest_pred)
        print(mse)
        mse = mean_squared_error(ytrain_pred,ytrain)
        print(mse)

33.448979997676524
19.326470203585725
```

```
In [39]: mse = mean squared error(ytest, ytest_pred)
```

```
In [47]: plt.scatter(ytrain ,ytrain_pred,c='blue',marker='o',label='Training data'),plt.scatter(ytest,ytest_pred ,c='r',marker='s')
        plt.title("True value vs Predicted value")
        plt.legend(loc='upper left')
        #plt.hlines(y=0,xmin=0,xmax=50)
        plt.plot()
        plt.show()
```



```
In [ ]:
```

```
In [ ]:
```

